



Computer Vision



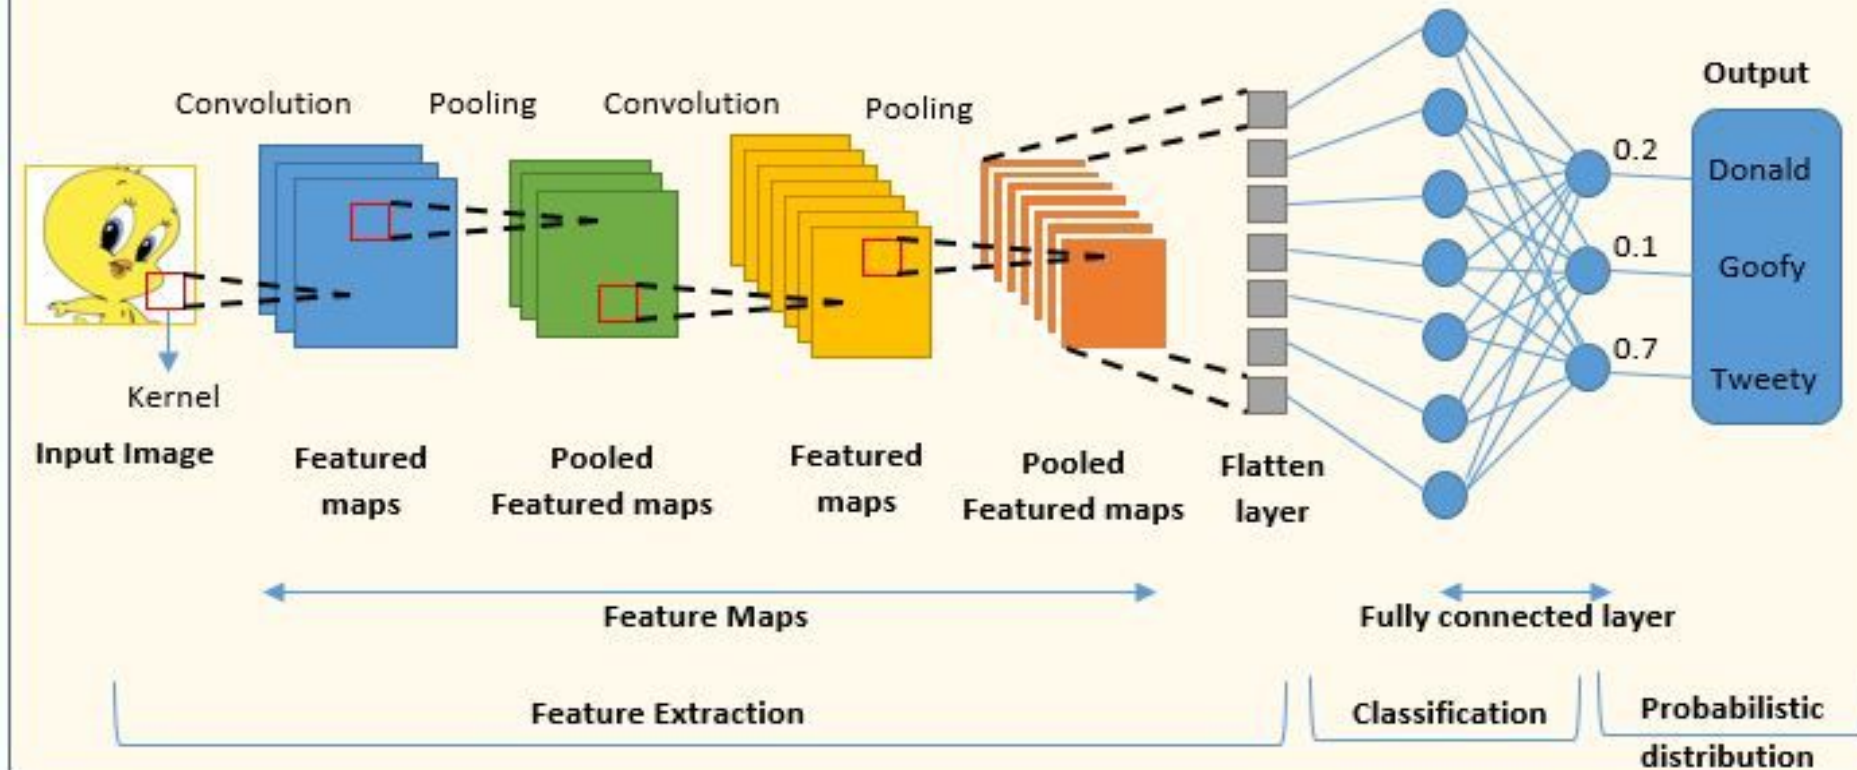
Dog Breed Classification

Goal: Build a model that will tell us the breed of our furry friends



Convolutional Neural Network

A Typical Convolutional Neural Network (CNN)



Issues With Our Data

- There were 20 and a half thousand images evenly distributed among the 120 breeds (not enough examples for the network!)
- Image data contained noise that needed filtering

Classification for 120 Breeds

We first started experimenting with our CNN on the entire dataset of 120 breeds.

This was a naive approach as the model produced very lackluster results.

Our accuracy was about 30% for the training dataset

Transfer Learning

Transfer Learning is a technique where you can build upon pre-trained models that have been tested on millions of samples and combine them with your own

This allows us to have a solid foundation for our own model and hopefully allow faster convergence.

We used the VGG16 package (a CNN) from Keras and the ImageNet training set of over 16 million different picture.

input_6	input:	[(None, None, None, 3)]
InputLayer	output:	[(None, None, None, 3)]

block1_conv1	input:	(None, None, None, 3)
Conv2D	output:	(None, None, None, 64)

block1_conv2	input:	(None, None, None, 64)
Conv2D	output:	(None, None, None, 64)

block1_pool	input:	(None, None, None, 64)
MaxPooling2D	output:	(None, None, None, 64)

block2_conv1	input:	(None, None, None, 64)
Conv2D	output:	(None, None, None, 128)

block2_conv2	input:	(None, None, None, 128)
Conv2D	output:	(None, None, None, 128)

block2_pool	input:	(None, None, None, 128)
MaxPooling2D	output:	(None, None, None, 128)

block3_conv1	input:	(None, None, None, 128)
Conv2D	output:	(None, None, None, 256)

block4_conv1	input:	(None, None, None, 256)
Conv2D	output:	(None, None, None, 512)

block4_conv2	input:	(None, None, None, 512)
Conv2D	output:	(None, None, None, 512)

block4_conv3	input:	(None, None, None, 512)
Conv2D	output:	(None, None, None, 512)

block4_pool	input:	(None, None, None, 512)
MaxPooling2D	output:	(None, None, None, 512)

block5_conv1	input:	(None, None, None, 512)
Conv2D	output:	(None, None, None, 512)

block5_conv2	input:	(None, None, None, 512)
Conv2D	output:	(None, None, None, 512)

block5_conv3	input:	(None, None, None, 512)
Conv2D	output:	(None, None, None, 512)

block5_pool	input:	(None, None, None, 512)
MaxPooling2D	output:	(None, None, None, 512)

global_average_pooling2d_5	input:	(None, None, None, 512)
GlobalAveragePooling2D	output:	(None, 512)

dense_20	input:	(None, 512)
Dense	output:	(None, 1024)

dense_21	input:	(None, 1024)
Dense	output:	(None, 1024)

dropout_5	input:	(None, 1024)
Dropout	output:	(None, 1024)

dense_22	input:	(None, 1024)
Dense	output:	(None, 512)

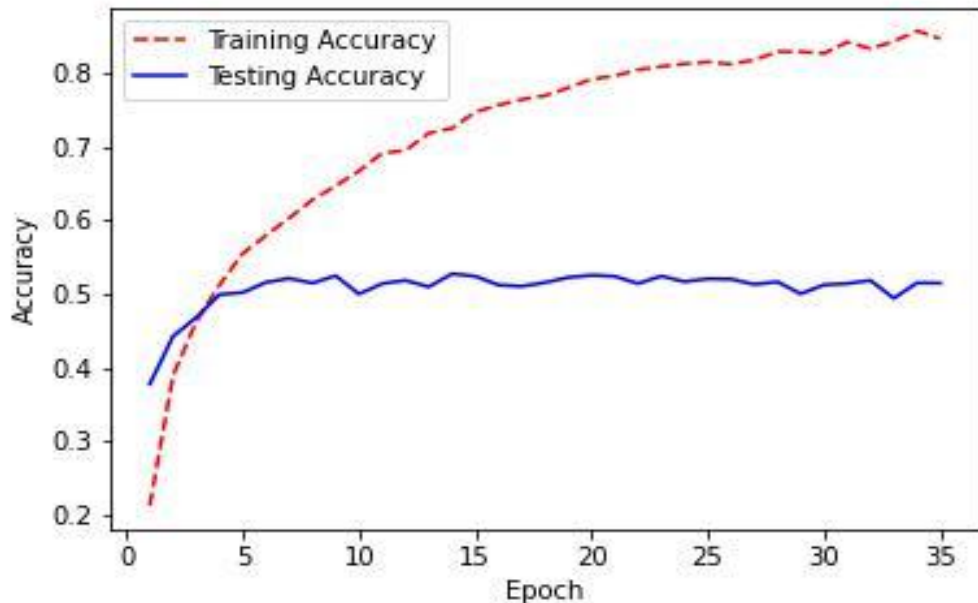
dense_23	input:	(None, 512)
Dense	output:	(None, 120)

120 Breed Model

Using the transfer learning technique we were able to get the following accuracies:

84.66% on the training set

51.5% on the testing set.



Here are some test predictions from the dataset.


Predicted: n02102973-Irish_water_spaniel
Actually: n02102973-Irish_water_spaniel




Predicted: n02093859-Kerry_blue_terrier
Actually: n02106382-Bouvier_des_Flandres




Predicted: n02102480-Sussex_spaniel
Actually: n02102480-Sussex_spaniel




Predicted: n02098413-Lhasa
Actually: n02098413-Lhasa



Predicted: n02099849-Chesapeake_Bay_retriever
Actually: n02099849-Chesapeake_Bay_retriever




Predicted: n02102318-cocker_spaniel
Actually: n02102318-cocker_spaniel




Predicted: n02111500-Great_Pyrenees
Actually: n02111500-Great_Pyrenees



Predicted: n02102480-Sussex_spaniel
Actually: n02102480-Sussex_spaniel

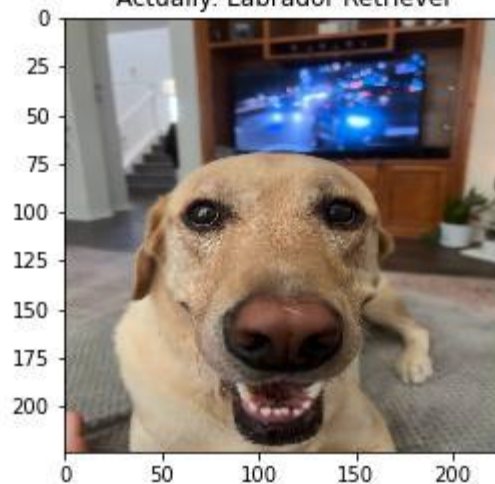


Predicted: n02112137-chow
Actually: n02093859-Kerry_blue_terrier



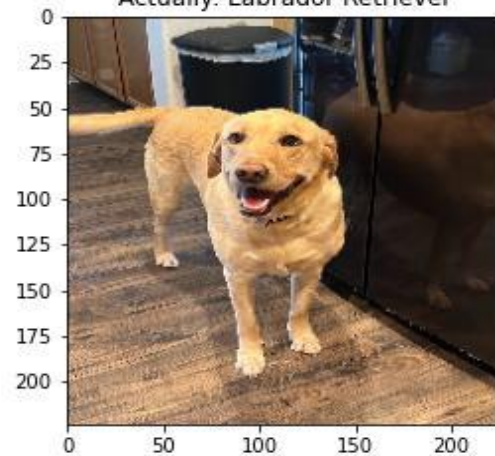
Predicted: ['n02092339-Weimaraner']

Actually: Labrador Retriever



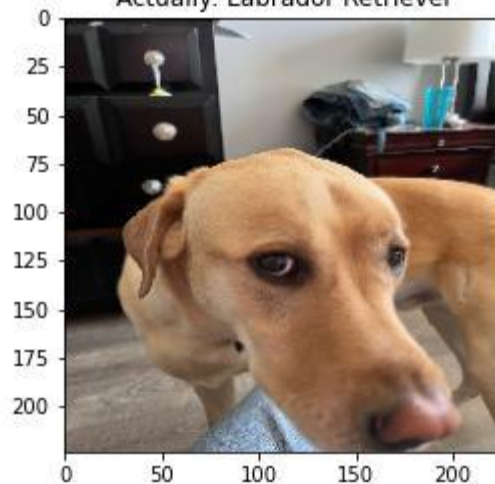
Predicted: ['n02099849-Chesapeake_Bay_retriever']

Actually: Labrador Retriever



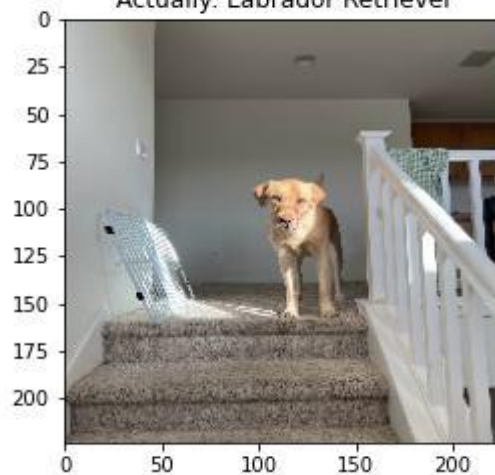
Predicted: ['n02092339-Weimaraner']

Actually: Labrador Retriever



Predicted: ['n02093428-American_Staffordshire_terrier']

Actually: Labrador Retriever



How to Improve?

There appears to be overfitting in our model, considering the discrepancy between the testing and training set.

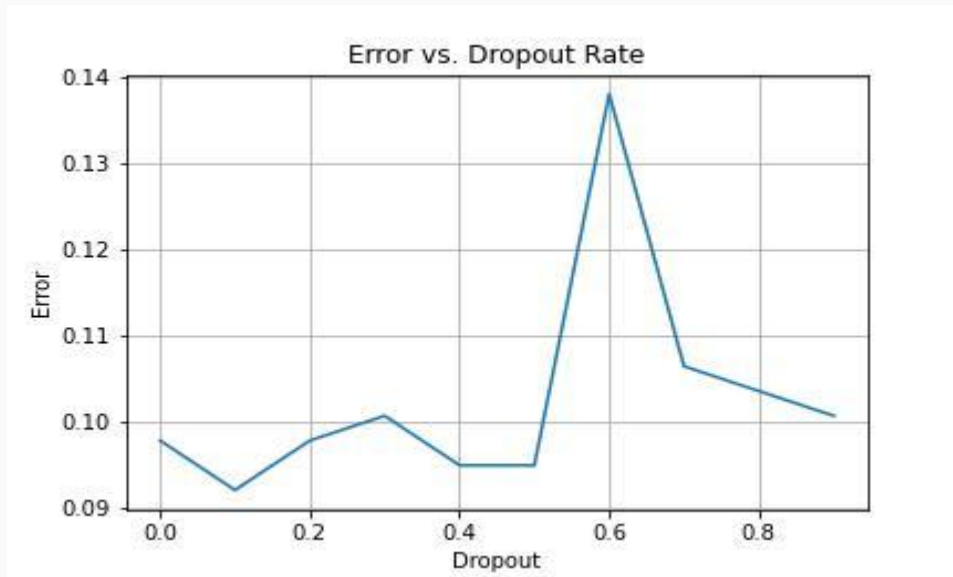
One possible way is to use the augmented data that John had created

The problem is, the data was 11 magnitudes larger than the original so it was not possible to run on our 120 breed classifier

10 Breed Classifier

Using the same methods as before, we used the augmented data and instead only ran the model using 10 different breed.

We experimented with different dropout rates and ended up with using the metric of 0.1

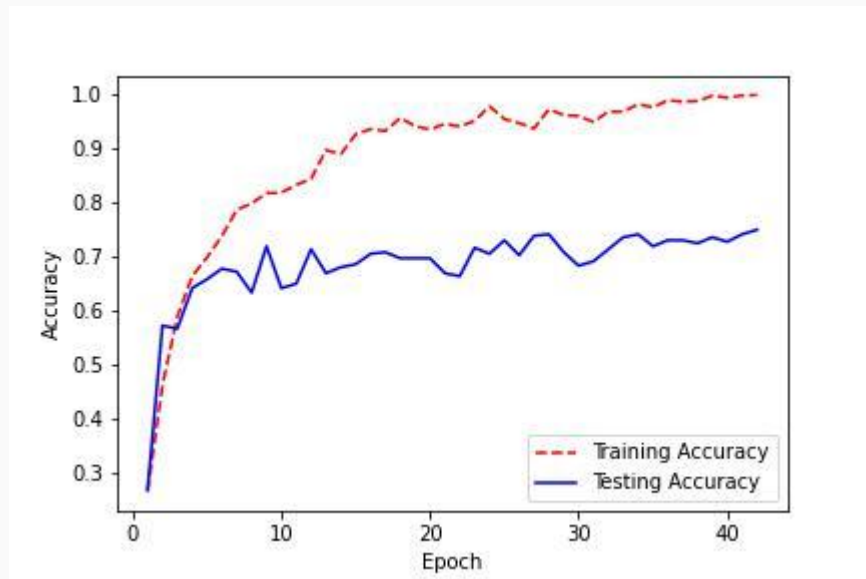


10 Breed Classifier


This model performed rather well and by using transfer learning and the augmented data we were able to get the following accuracies:

100% on the training set

75.1% on the testing set.



Some predictions using the 10 breed classifier



Summary

If we were able to use more memory for our model, we could certainly get a lower error rate for our 120 breed classifier. It is simply unfeasible to attempt it without additional GPU processing power.

A 25% error rate for a problem as complex as this is not so bad. Classifying dog breeds in pictures from real world data is a monumental task. Still, there are even more methods we could use to improve our model such as filtering for the images and changing the parameters of the Neural Network.

Questions?